QUANTIFYING AND IMPROVING THE SECURITY OF BLOCKCHAIN SYSTEMS

A Dissertation Outline Presented

by

A. PINAR OZISIK

Submitted to the Graduate School of the
University of Massachusetts Amherst in partial fulfillment
of the requirements for the degree of

DOCTOR OF PHILOSOPHY

September 2019

College of Information and Computer Sciences

QUANTIFYING AND IMPROVING THE SECURITY OF BLOCKCHAIN SYSTEMS

A Dissertation Outline Presented

by

A. PINAR OZISIK

Approved as to style and content by:

_____

Brian N. Levine, Chair

_____

Philip S. Thomas, Member

_____

Yuriy Brun, Member

_____

Nikunj Kapadia, Member

_____

James Allan, Chair
College of Information and Computer Sciences

ABSTRACT

QUANTIFYING AND IMPROVING THE SECURITY OF BLOCKCHAIN SYSTEMS

A. PINAR OZISIK

B.Sc., BRANDEIS UNIVERSITY

M.S., UNIVERSITY OF MASSACHUSETTS

Directed by: Professor Brian N. Levine

In this thesis, I analyze and improve the security and performance of blockchain systems across three primary themes. In the first theme, I analyze blockchain algorithms for setting block discovery difficulty. Unfortunately, churn in mining power can cause uneven inter-block delays when the difficulty is not set accurately. Mining power can change due to many reasons, including the miners' allocation of hardware and swings in the exchange rate of a currency. For example, Bitcoin Cash has seen enormous variance in mining power since its creation and the existing algorithm for difficulty did not easily converge. I propose two alternatives to accurately update difficulty: one that solely uses information that is currently available in blockchain networks, and another based on status reports regularly broadcast from some or all miners of their partial proof-of-work (POW). Status reports can also be used for emergency difficulty adjustment, an algorithm the network resorts to when a block takes unusually long to discover.

Status reports add overhead into networks because they require the broadcast of additional information. In a second theme, I introduce a novel method of interactive set reconciliation for the distribution of status reports in order to reduce traffic. Even without status reports, this protocol works for the efficient distribution of blocks. The approach, called Graphene, couples a Bloom filter with an IBLT. Then I evaluate performance analytically and show that Graphene blocks are always smaller and therefore network performance is improved.

In the third theme, I analyze the practical feasibility of double-spend and selfish mining attacks on blockchain systems. The hash rate of miners is the primary quantitative factor that determines the security of any POW based blockchain consensus algorithm. Most analyses generally assume that the hash rate of honest and malicious miners is known. However, I show that hash rate estimation is difficult and introduces high variance. Therefore, I argue that these double-spend and selfish mining attacks are difficult to carry out with high precision, and use reinforcement learning techniques to realistically evaluate these attacks when an attacker does not have full knowledge of the networks mining power.

TABLE OF CONTENTS

LIST OF TABLES

LIST OF FIGURES

INTRODUCTION

Contributions

The following is a summary of the contributions in each chapter of this proposal.

1.

Collaborators

All research activities are conducted under the supervision of Brian Levine. Preliminary work for Chapter 2 was completed in collaboration with George Bissias and Brian Levine; that for Chapter 3 was completed in collaboration with Gavin Andresen, George Bissias, Amir Houmansadr and Brian Levine; and that for Chapter 4 was completed in collaboration with Philip Thomas and Brian Levine.

CHAPTER 1

OVERVIEW OF BLOCKCHAIN SYSTEMS

In this chapter, we describe the basic operation of blockchains using Bitcoin as an example. Other cryptocurrencies such as Litecoin and Ethereum operate similarly with minor differences.

1.1 Basic Operation

Accounts. A Bitcoin is a unit of currency, which is fungible, divisible (up to eight decimal places), and recombinable. It is measured as a balance across multiple accounts, which are themselves manifested in *addresses*.[1] Each address comprises a stored asymmetric cryptographic key and an associated balance of Bitcoin. The public portions of an address are the public key and the balance of coin. When an address is involved in a *transaction* with one or more other addresses, Bitcoins are transferred among them. Addresses are explicitly pseudonyms, and not tied to a particular individual; further, empty addresses can be created at no cost beyond generating an asymmetric key pair.

Adding to the blockchain. To be added to the *blockchain*, transactions are broadcast by users on Bitcoin's peer-to-peer (p2p) network. A set of *miners* on the p2p network verify that each transaction is signed correctly, does not conflict with a previous transaction, does not move more coin than is contained in the address, and other functions. Each miner independently agglomerates a set of valid transactions

---

[1]Internally, Bitcoins exist only as "unspent transaction outputs" (UTXO), but users of the system think of them as balances in addresses.

into a candidate *block* and attempts to solve a predefined cryptographic puzzle as proof-of-work (POW), which involves data from the candidate block and a specific *prior block*. The new transactions are only valid if they do not conflict with the set of transactions that are contained in all blocks that are direct ancestors.

The first miner to solve the problem broadcasts his solution to the network, and by virtue of the solution, is able to add the block to the ever-growing blockchain as a child of the prior block. The miners then start over, using the newly appended blockchain and the set of remaining transactions. The miners' incentive for *discovering* a new block is a reward of coins, called the *coinbase*, consisting of a predetermined *block reward* (currently worth 12.5 Bitcoins) and fees from transactions included in the block.

In Bitcoin, the POW computation is dynamically calibrated to take approximately ten minutes per block. When transactions appear in a block, they are *confirmed*, and each subsequent block provides additional confirmation. To announce a new block, a miner lists all transactions contained in the new block along with a header that contains an easily-verifiable POW solution. When a node or miner receives a new block, he validates each transaction in the block and the POW.

Notably, if there is a fork on the chain, honest miners always select the prior block as the last block containing the largest amount of POW. However, due to propagation delays in the network, it is possible for the miners to receive competing (but valid) block announcements, which bifurcates the chain, until one of the two forks is appended to first. It is also possible and valid for a miner to receive a set of blocks that retroactively rewrites many blocks; doing so is a demonstration of computational work that miners accept despite the age or depth[2] of a rewritten block.

---

[2]The *depth* of a block refers to the number of blocks that follow it; the *height* of a block is the number of blocks that precede it.

Any entity can elect to be a miner for Bitcoin, and there is no centralized party from whom to seek approval for mining. If all miners were to simply vote on which block should be appended to the main chain, then the mining process would fall vulnerable to a Sybil attack [7]. The POW puzzle addresses this problem by performing a kind of decentralized leader election: the miner that solves the puzzle can decide which block to append to the chain.

Full nodes. *Full nodes* are peers in the network that do not mine, but do generate, validate, and propagate transactions and blocks to other nodes including miners. Consumers (i.e., those who purchase goods or services) typically have no need to process and validate all transactions, so they can instead operate *simple payment verification* (SPV) nodes that process, store, and transmit data involving only addresses-of-interest, which are typically addresses they control, make payments to, or receive payments from. SPV nodes rely on full nodes to relay transactions-of-interest.

Bitcoin transaction consistency. The main goal of the p2p network is to provide a consistent view of blocks and unconfirmed transactions across all network peers. Each peer maintains a local snapshot of the transactions in a memory pool dubbed the *mempool*. Blocks consist of a list of transactions that have already (almost always) been broadcast to miners and full nodes in the network.

CHAPTER 2

DIFFICULTY ESTIMATION

2.1   POW in Blockchain Systems

Bitcoin uses a simple POW algorithm based on cryptographic hashing, proposed earlier by Douceur [7]. Specifically, miners apply a 256-bit cryptographic hash algorithm [14] to an 80-byte *block header*, and the puzzle is solved if the resulting value is less than a known *target*, $0 < t < 2^{256}$. The header in Bitcoin consists of the Merkle root of the set of transactions, a timestamp, the target (stored as $2^{224}/t$), a *nonce*, and the hash of the prior block's header. If the hash is not less than the target, then a new nonce is selected to generate a new hash (the Merkle root can be adjusted as well). This process repeats until some miner finds a solution.

Each time a nonce is selected and the block header is hashed, the miner is sampling a value from a discrete uniform distribution with range $[0, 2^{256}-1]$. The probability of solving the POW and discovering a block is the cumulative probability of selecting a value from $[0, t]$, which is $t/2^{256}$. Hence, in expectation, the number of samples needed to discover a block is $2^{256}/t$. Bitcoin adjusts the target so that on average it takes about 600 seconds to find a block. Typically, the target is described for convenience as a *difficulty*, defined to be $D = 2^{224}/t$. Bitcoin's difficulty is set once every two weeks.

Ethereum. Ethereum operates very similarly to Bitcoin. Miners solve a POW problem that is more complicated than Bitcoin in an attempt to disadvantage miners with custom ASICs. However, in the end, a miner still compares a hash value to the target. Specifically, the number of values in the block header is larger, resulting in a

508-byte header. It's not the hash of the header that is compared against the target, but the hash resulting from an Ethereum-specific algorithm called ETHASH [9], for which the hash of the block header is the primary input. In the end, the POW hash value is a sample from a discrete uniform distribution with range $[0, 2^{256} - 1]$, and the probability of block discovery is $t/2^{256}$.

A major difference of Ethereum is that the target is set such that the expected time between blocks is 15 seconds. This setting results in quicker confirmation times, but as a result, the probability that two miners announce blocks within the propagation time of a block announcement is much higher. Therefore, there are many abandoned forks in the chain. Ethereum uses a modified version of the GHOST [19] protocol for selecting the main fork of the blockchain: the main chain follows the block at each level with the most POW on its subtree. These differences do not affect the application of our algorithms; in fact, the presence of ommers is additional data which improves our estimates.

2.2 Problem Statement

Developers have resorted to ad-hoc methods for updating the difficulty in many blockchain systems. So far, there has been no previous work on analyzing the efficiency and correctness of these methods. In fact, because some blockchain systems do not accurately update difficulty, networks see enormous variance in inter-block delay. If mining power were constant in these networks, then difficulty could be kept constant. However, as seen with many blockchain systems, including a recent example with Bitcoin Cash, mining power fluctuates within these networks, requiring a readjustment of network parameters, the most important being difficulty. Therefore, in this chapter, I use the most prominent cryptocurrency, Bitcoin, as a testbed for analyzing difficulty, and then propose alternative methods that could potentially increase efficiency.

2.3 Preliminary Work

2.3.1 Analysis of Bitcoin's Difficulty Estimation

Algorithm for setting difficulty. Bitcoin's networks parameters are set such that a block is discovered every 10 minutes. Initially, difficulty starts at 1, and then for every 2016 blocks that are found, the timestamps of the blocks are compared to find out how much time it took to find 2016 blocks. Let t denote the time in minutes it took to find 2016 blocks. Because the network is configured such that 2016 blocks must take 2 weeks (20160 minutes), the old difficulty is multiplied by $20160/t$. If the correction factor is greater than 4 or less than $1/4$, then 4 or $1/4$ are used, respectively, to prevent the change from being too abrupt.

Bitcoin's target and difficulty are related to each other as follows: target $=$ targetmax $/$ difficulty $= 2^{224}/D_i$, where D_i is the ith time the difficulty is set [6]. The difficulty and target are inversely proportional: when difficulty increases, the target decreases. Therefore, a smaller target makes block creation more difficult, and as the difficulty goes up, so does the expected time needed to create a block.

Analysis of difficulty adjustment. Let D_i denote the ith time the difficulty is set, and X_k denote the number of minutes it took to generate the kth block after D_i is set. Then we have

$$D_i = D_i \frac{10n}{\sum_{k=1}^{n} X_k}. \tag{2.1}$$

Let D be a sequence generated by the last equation where the first element of the sequence is $D_0 = 1$. Mining is an example of a Poisson process because, under constant mining power, blocks are mined continuously and independently at a constant average rate. Therefore, $X_k \sim \text{Exp}(\beta)$ with $\beta = 1/\lambda$, assuming that the miner hash rate stays constant for the 2-week period after D_i is set. In this parametrization of the exponential, β represents the survival parameter, and hence, the ratio describing

the *expected time* it takes for one block to arrive. For example, ideally in Bitcoin $\beta = 1/10$ minutes and in Ethereum $\beta = 1/15$ seconds. The difficulty for the $i + 1$st time given D_i and $X_1, \ldots, X_n$ is

$$D_{i+1} = 10nD_i \left(\frac{1}{\sum_{k=1}^{n} X_k} \right). \tag{2.2}$$

The relationship between hash rate and β. Given difficulty D_i, the expected number of hashes, h, needed to meet the target for a block is

$$\mathbb{E}[h] = \frac{2^{256} - 1}{T_i} = \frac{2^{256} - 1}{2^{224}/D_i} = \frac{D_i(2^{256} - 1)}{2^{224}}, \tag{2.3}$$

where T_i is the target set for the ith time. Note that $\mathbb{E}[h]$ describes the *total* number of expected hashes needed to discover a block, and I have observations regarding the *time* it takes to generate a block. Let r be the hash rate of the network in minutes (or the number of hashes per time unit), and $X = X_1, \ldots, X_n$, where $X \sim \text{Exp}(\beta)$, with $\beta = 1/\lambda$. λr is the expected number of hashes each time a block is created.

$$\mathbb{E}[h] = r\lambda = r\frac{1}{\beta} \tag{2.4}$$

$$r = \mathbb{E}[h]\beta. \tag{2.5}$$

Adjusting difficulty correctly. For Bitcoin, where the network is expected to solve a block every 10 minutes, we can adjust the target for the $i + 1$th time as follows

$$\frac{(2^{256} - 1)}{T_{i+1}} = 10r \tag{2.6}$$

$$\frac{(2^{256} - 1)}{T_{i+1}} = \frac{10(2^{256} - 1)\beta}{T_i} \tag{2.7}$$

$$T_{i+1} = \frac{T_i}{10\beta}. \tag{2.8}$$

8

Therefore, by rearranging and substituting T_i with its definition using D_i, we can adjust the difficulty for the $i + 1$th time as follows

$$D_{i+1} = \frac{2^{224}}{T_{i+1}} \tag{2.9}$$

$$= 10\beta D_i. \tag{2.10}$$

Variance of difficulty. Additionally, we can also talk about the variance of a term in sequence D, given its preceding term and the new data we see.

$$\text{Var}(D_{i+1}|D_i, X_1, \ldots, X_n) = \text{Var}\left(10nD_i\frac{1}{Y}\right) \tag{2.11}$$

$$= (10nD_i)^2\text{Var}\left(\frac{1}{Y}\right) \tag{2.12}$$

$$= (10nD_i)^2\left(\frac{1}{\beta^2(n-1)^2(n-2)}\right) \tag{2.13}$$

$$= \frac{(10nD_i)^2}{\beta^2(n-1)^2(n-2)}. \tag{2.14}$$

Bias of difficulty.

$$\text{bias}(D_{i+1}|D_i, X_1, \ldots, X_n) = \mathbb{E}[D_{i+1}|D_i, X_1, \ldots, X_n] - D_{i+1} \tag{2.15}$$

$$= \frac{10nD_i}{\beta(n-1)} - 10\beta D_i. \tag{2.16}$$

Mean squared error (MSE) of difficulty.

$$\text{MSE}(D_{i+1}|D_i, X_1, \ldots, X_n) = \text{bias}(D_{i+1}|D_i, X_1, \ldots, X_n)^2 + \text{Var}(D_{i+1}|D_i, X_1, \ldots, X_n) \tag{2.17}$$

$$= \left(\frac{10nD_i}{\beta(n-1)} - 10\beta D_i\right)^2 + \frac{(10nD_i)^2}{\beta^2(n-1)^2(n-2)}. \tag{2.18}$$

2.3.2 Alternative Time-Based Estimation

Estimator for β, the expected inter-arrival time between blocks. Let $X = X_1, \ldots, X_n$ denote the inter-arrival time between $n + 1$ consecutive blocks on the blockchain. Given a consecutive sequence of $n + 1$ blocks, n inter-arrival times can be computed by subtracting the timestamp of each block from that of its preceding block. Note that $X \sim \text{Exp}(\beta)$ with $\beta = 1/\lambda$, similar to the definition in the previous section. It it well known that the unbiased MLE estimator for β is

$$\hat{\beta} = \frac{\sum_{k=1}^{n} X_i}{n}. \tag{2.19}$$

Adjusting difficulty. Using our estimation of β, we can adjust the difficulty for the $i + 1$th time as follows

$$D_{i+1} = 10 D_i \hat{\beta}. \tag{2.20}$$

Variance of New Difficulty.

$$\text{Var}(D_{i+1}|D_i, X_1, \ldots, X_n) = \text{Var}(10 D_i \hat{\beta}) \tag{2.21}$$
$$= (10 D_i)^2 \text{Var}(\hat{\beta}) \tag{2.22}$$
$$= \frac{(10 D_i \beta)^2}{n}. \tag{2.23}$$

Bias of New Difficulty.

$$\text{bias}(D_{i+1}|D_i, X_1, ..., X_n) = \mathbb{E}[D_{i+1}|D_i, X_1, ..., X_n] - D_{i+1} \tag{2.24}$$
$$= 10 D_i \beta - 10 D_i \beta \tag{2.25}$$
$$= 0. \tag{2.26}$$

MSE of New Difficulty.

$$\text{MSE}(D_{i+1}|D_i, X_1, ..., X_n) = \text{bias}(D_{i+1}|D_i, X_1, ..., X_n)^2 + \text{Var}(D_{i+1}|D_i, X_1, ..., X_n)$$

$$\tag{2.27}$$

$$= \frac{(10D_i\beta)^2}{n}. \tag{2.28}$$

2.3.3 Comparison of Estimators

Note that our alternative estimator has zero bias compared to Bitcoin's original estimator. Additionally, under the appropriate constraints, variance and MSE is also significantly lower.

Variance.

$$\frac{(10D_i\beta)^2}{n} \leq \frac{(10nD_i)^2}{\beta^2(n-1)^2(n-2)} \tag{2.29}$$

$$\frac{\beta^2}{n} \leq \frac{n^2}{\beta^2(n-1)^2(n-2)} \tag{2.30}$$

Our estimator (LHS) has lower variance than the original estimator (RHS) for $n > 2$ and $0 < \beta < 1$.

MSE.

$$\frac{(10D_i\beta)^2}{n} \leq \left(\frac{10nD_i}{\beta(n-1)} - 10\beta D_i\right)^2 + \frac{(10nD_i)^2}{\beta^2(n-1)^2(n-2)} \tag{2.31}$$

Our estimator (LHS) has lower variance than the original estimator (RHS) for $n > 2$, $0 < \beta < 1$ and $D_i \in \mathcal{R}$. [Add figure for Bitcoin, Litecoin, Ethereum, Bitcoin Cash, Ripple. After 2016 blocks generated, Litecoin will adjust difficulty to estimated difficulty in order to keep the block generation time at 150 seconds. bitcoin cash has a window of only 144 blocks. and uses a sliding window. Ethereum uses 1 block, can't find ripple]

2.4 Proposed Work

We were able to show that our estimator performs better than Bitcoin's ad-hoc method. To extend this chapter, I propose to answer a few more key questions:

1. Analyze different attacks against our estimator. What happens when the timestamps on the blocks are reported inaccurately by an attacker? How much error can an attacker introduce?

2. How often should difficulty be adjusted? Should a sliding window or non-overlapping window of blocks be used?

3. Given a required error rate, what is the shortest window of time (or number of blocks) needed to estimate difficulty correctly?

4. Given that mining power is changing, how quickly can either estimator adapt to change?

5. Attackers can easily lie about the timestamps on the blocks with our new estimator. Therefore, an alternative estimator I plan to develop uses status reports, which is a block header except that the POW does not satisfy the current target. Instead, the minimum hash value in the report represents the hash found since the last block broadcast on the chain. To be clear, each status report does not directly report the minimum hash value; instead, reports are of the input values to the POW algorithm. Because attackers can't lie about their POW, an estimator based on the minimum hash value is safer and can be used to adjust the emergency difficulty when the network fails to produce a block for an unusually long time. What is the variance and MSE of the status report based estimator?

6. What are the advantages and disadvantages of status reports compared to the time-based estimator?

7. What are the Chernoff bounds associated with these estimators?

CHAPTER 3

GRAPHENE: EFFICIENT BLOCK ANNOUNCEMENTS

3.1 Background

In this section, we describe the signaling mechanism behind blockchain systems, explain the operation of IBLTs and summarize related work.

3.1.1 Topology and Signaling in Blockchain Systems

Bitcoin propagates new transaction and block announcements by flooding throughout a p2p random graph of full nodes and miners. Each peer in the graph requests direct connections to 8 other peers, and accepts requests for connections from up to 117 other peers. A peer will offer a newly created transaction to each neighbor via an `inv` message, which reports the hash of the transaction content as its ID. If a peer does not already possess the transaction, it will request it using a `getdata` message. Blocks are handled similarly: `inv` messages describe a block by its ID, which is created from the hash of the block's contents. Upon receiving the `inv`, peers will request the block if they do not already have it. Hence, in today's topology, `inv` messages cross every edge in the random graph once, while the actual transaction and block data typically propagate along only a spanning tree of the graph (more edges will be traversed if there are propagation delays). For convenience, we refer to the set of (unconfirmed) transaction IDs that a peer knows about as the *IDpool*. Actual transaction contents are placed in the mempool.

3.1.2 Related Work

The main limitation we are addressing with Graphene is the inefficiency of blockchain systems in disseminating block data. A block announcement must be validated using the transaction content comprising the block. However, it is likely that the majority of the peers have already received these transactions, and they only need to discern them from those in their mempool. Additionally, in Section 2.4, we propose using status reports, blocks announcement that do not satisfy the difficulty requirement of the network, for the emergency difficulty algorithm. Therefore, on top standard blocks, we are proposing to add more network traffic, requiring an even more efficient method of block propagation.

In principle, a block announcement needs to include only the IDs of those transactions, and accordingly, Corallo's *Compact Block* design [5] — which has been recently deployed — significantly reduces block size by including a transaction ID list at the cost of increasing coordination to 3 roundtrip times. *Xtreme Thinblocks* [21], an alternative protocol, works similarly to Compact Blocks but has greater data overhead. Specifically, after receiving an `inv` for a block, the receiver creates a Bloom filter of her mempool with FPR $f = 1/n$, where n is the number of transactions in the block. The sender then sends a *thinblock transaction* that contains block header information, all transaction IDs in the block and any transactions that do not pass through the Bloom filter, enabling the receiver to recreate the block. As a result, Xtreme Thinblocks are larger than Compact Blocks but require just 2 roundtrip times. Relatedly, the community has discussed in forums the use of IBLTs (alone) for reducing block announcements [1,17], but these schemes have not been formally evaluated and are less efficient than our approach. Our novel method, which we prove and demonstrate is smaller than all of these recent works, requires just 2 roundtrip times for coordination.

3.1.3 Overview of IBLTs

We make use of Invertible Bloom Lookup Tables (IBLTs) [12], which is an efficient data structure for *set reconciliation* between two peers. Like Bloom filters [3], IBLTs allow two parties to determine, with high probability, which values from a set they share in common. But unlike Bloom filters, IBLTs enable the recovery of any missing values, which are assumed to be of fixed size and encoded as binary strings. Key-value pairs can be inserted, retrieved and deleted like an ordinary hash table. An IBLT consists of m entries, each storing a `count`, a `keySum`, and a `valueSum`, all initialized to zero.

A new value v is inserted into location $i = h(v)$ based on the hash of its value such that $i < m$. At entry i, all three fields are incremented or xored. In particular, standard addition is used for the `count` field, but `xor` is used to add to the `keySum` and `valueSum` fields. An item can be deleted similarly: at the correct entry, `count` is subtracted by 1, and the `valueSum` and `keySum` fields are `xor`'ed. When `count` $\equiv 1$ the `valueSum` field contains the actual value of the sole item remaining in the cell. (The purpose of the `keySum` field is to support a `GET()` operation for a given key: that is, if `count` $\equiv 1$ and `keySum` $\equiv h(v)$, then `valueSum` $\equiv v$.) IBLTs use $k > 1$ hash functions to store each value in k entries, which we collectively call a value's *entry set*. If table space is sufficient, then with high probability for at least one of the k entries, `count` $\equiv 1$.

Suppose that two peers each have a list of values, V and V', respectively, such that the difference is expected to be small. The first peer constructs an IBLT L (with m entries) from V. The second peer constructs V' from L' (also having m entries). Eppstein et al. [8] showed that a cell-by-cell difference operator can be used to efficiently compute the symmetric difference $L \triangle L'$. For each pair of fields (f, f'), at each entry in L and L', we compute either $f \oplus f'$ or $f - f'$ depending on the field type. When $|\texttt{count}| \equiv 1$ at any entry, the corresponding value can be

recovered. Peers proceed by removing the recoverable key-value pair from all entries in the value's entry set. This process will generally produce new recoverable entries, and continues until nothing is recoverable.

3.2 Graphene

In this section, we detail *Graphene*, where a receiver learns the set of specific transaction IDs that are contained in a (pending or confirmed) block containing n transactions. Unlike other approaches, Graphene never sends an explicit list of transaction IDs, instead it sends a small Bloom filter and a very small IBLT.

PROTOCOL 1: Graphene

1: **Sender:** Sends `inv` for a block.

2: **Receiver:** Requests unknown block; includes count of txns in her IDpool, m.

3: **Sender:** Sends Bloom filter $\mathcal{S}$ and IBLT $\mathcal{I}$ (each created from the set of n txn IDs in the block) and essential Bitcoin header fields. The FPR of the filter is $f = a/(m - n)$, where $a = n/(c\tau)$.

4: **Receiver:** Creates IBLT $\mathcal{I}'$ from the txn IDs that pass through $\mathcal{S}$. She decodes the *subtraction* [8] of the two blocks, $\mathcal{I} \bigtriangleup \mathcal{I}'$.

The intuition behind Graphene is as follows. The sender creates an IBLT $\mathcal{I}$ from the set of transaction (txn) IDs in the block. To help the receiver create the same (or similar) IBLT, he also creates a Bloom filter $\mathcal{S}$ of the transaction IDs in the block. The receiver uses $\mathcal{S}$ to filter out transaction IDs from her IDpool and creates her own IBLT $\mathcal{I}'$. She then attempts to use $\mathcal{I}'$ to *decode* $\mathcal{I}$, which, if successful, will yield the transaction IDs comprising the block. The number of transactions that falsely appear to be in $\mathcal{S}$, and therefore are wrongly added to $\mathcal{I}'$, is determined by a parameter controlled by the sender. Using this parameter, he can create $\mathcal{I}$ such that it will decode with very high probability.

17

A Bloom filter is an array of x bits representing y items. Initially, the x bits are cleared. Whenever an item is added to the filter, k bits, selected using k hash functions, in the bit-array are set. The number of bits required by the filter is $x = -y\ln(f)/\ln^2(2)$, where f is the intended false positive rate (FPR). For Graphene, we set $f = a/(m - n)$, where a is the expected difference between $\mathcal{I}$ and $\mathcal{I}'$. Since the Bloom filter contains n entries, and we need to convert to bytes, its size is

$$\frac{-\ln(\frac{a}{m-n})}{\ln^2(2)}\frac{1}{8}.$$

It is also the case that a is the primary parameter of the IBLT size. IBLT $\mathcal{I}$ can be decoded by IBLT $\mathcal{I}'$ with very high probability if the number of cells in $\mathcal{I}$ is d-times the expected symmetric difference between the list of entries in $\mathcal{I}$ and the list of entries in $\mathcal{I}'$. In our case, the expected difference is a, and we set $d = 1.5$ (see Eppstein et al. [8], which explores settings of d). Each cell in an IBLT has a `count`, `keySum` and `valueSum`. (It can also have a key, but we have no need for a key). For us, the count field is 2 bytes, the keySum is 4 bytes, and the valueSum is the last 5 bytes of the transaction ID (which is sufficient to prevent collisions). In sum, the size of the IBLT with a symmetric difference of a entries is $1.5(2 + 4 + 5)a = 16.5a$ bytes. Thus, the total cost in bytes, T, for the Bloom filter and IBLT are given by

$$T(a) = n\frac{-\ln(f)}{c} + a\tau = n\frac{-\ln(\frac{a}{m-\mu})}{c} + a\tau,$$

where all Bloom filter constants are grouped together as $c = 8\ln^2(2)$, and we let the overhead on IBLT entries be the constant $\tau = 16.5$.

To set the Bloom filter as small as possible, we must ensure that the FPR of the filter is as high as permitted. If we assume that all `inv` messages are sent ahead of a block, we know that the receiver already has all of the transactions in the block in her IDpool (they need not be in her mempool). Thus, $\mu = n$; i.e., we allow for a of

$m - n$ transactions to become false positives, since all transactions in the block are already guaranteed to pass through the filter. It follows that

$$T(a) = n\frac{-\ln(\frac{a}{m-n})}{c} + a\tau. \tag{3.1}$$

Taking the derivative w.r.t. a, Eq. 3.1 is minimized[1] when when $a = n/(c\tau)$.

Due to the randomized nature of an IBLT, there is a non-zero chance that it will fail to decode. In that case, the sender resends the IBLT with double the number of cells (which is still very small). In our simulations, presented in the next section, this doubling was sufficient for the incredibly few IBLTs that failed.

PROTOCOL 2: CompactBlocks

1:	Sender:	Sends `inv` for a block that has n txns.
2:	Receiver:	If block is not in mempool, requests compact block.
3:	Sender:	Sends the block header information, all txn IDs in the block and any full txns he predicts the sender hasn't received yet.
4:	Receiver:	Recreates the block and requests missing txns if there exist any.

3.3   Comparison to Compact Blocks

In this section, we mathematically show the efficiency of Graphene to Compact Blocks and present our simulation results comparing the two protocols.

---

[1]Actual implementations of Bloom filters and IBLTs involve several (non-continuous) ceiling functions such that we can re-write:

$$T(a) = \left(\lceil\ln(\frac{m-n}{a})\rceil \left\lceil \frac{n\ln(\frac{m-n}{a})}{\lceil\ln(\frac{m-n}{a})\rceil \ln^2(2)} \right\rceil\right) \frac{1}{8} + \lceil a\rceil\tau. \tag{3.2}$$

The optimal value of Eq. 3.2 can be found with a simple brute force loop. We compared the value of a picked by using $a = n/(c\tau)$ to the cost for that a from Eq. 3.2, for valid combinations of $50 \leq n \leq 2000$ and $50 \leq m \leq 10000$. We found that it is always within 37% of the cost of the optimal value from Eq. 3.2, with a median difference of 16%. In practice, a for-loop brute-force search for the lowest value of a is almost no cost to perform, and we do so in our simulations.

3.3.1 Mathematical Analysis

Compact Blocks [5] is to our knowledge the best-performing related work. It has several modes of operation, and we examined the *Low Bandwidth Relaying* mode due to its bandwidth efficiency, which operates as follows. After fully validating a new block, the sender sends an `inv`, for which the receiver sends a `getdata` message if she doesn't have the block. The sender then sends a `compact block` that contains block header information, all transaction IDs (shortened to 5 bytes) in the block, and any transactions that he predicts the receiver does not have (e.g., the coinbase). If the receiver still has missing transactions, she requests them via an `inv` message. Protocol 2 outlines this mode of Compact Blocks. The main difference between Graphene and Compact Blocks is that instead of sending a Bloom filter and an IBLT, the sender sends block header information and all shortened transaction IDs to the receiver.

A detailed example of how to calculate the size of each scheme is below; but we can state more generally the following result. For a block of n transactions, Compact Blocks costs $5n$ bytes. For both protocols, the receiver needs the `inv` messages for the set of transactions in the block before the sender can send it. Therefore, we expect the size of the IDpool of the receiver, m, to be constrained such that $m \geq n$. Assuming that $m > 0$ and $n > 0$, the following inequality must hold for Graphene to outperform Compact Blocks:

$$n\frac{-\ln(\frac{a}{m-n})}{c} + a\tau \quad < \quad 5n \qquad (3.3)$$

$$n \quad > \quad \frac{m}{1287670}. \qquad (3.4)$$

In other words, Graphene is strictly more efficient than Compact Blocks *unless* the set of unconfirmed transactions held by peers is 1,287,670 times larger than the block size (e.g., over 22 billion unconfirmed transactions for the current block size.) Finally, we note that Xtreme Thinblocks [21] are strictly larger than Compact Blocks since

they contain all IDs and a Bloom filter, and therefore Graphene performs strictly better than Xtreme Thinblocks as well.

Example. A receiver with an IDpool of $m = 4000$ transactions makes a request for a new block that has $n = 2000$ transactions. The value of a that minimizes the cost is $a = n/(c\tau) = 31.5$. The sender creates a Bloom filter $\mathcal{S}$ with $f = a/(m - n) = 31.5/2000 = 0.01577$, with total size of $2000 \times -ln(0.01577)/c = 2.1$ KB. The sender also creates an IBLT with a cells, totaling $16.5a = 521B$. In sum, a total of $2160B + 521B = 2.6$ KB bytes are sent. The receiver creates an IBLT of the same size, and using the technique introduced in Eppstein et al. [8], the receiver subtracts one IBLT from the other before decoding. In comparison, for a block of n transactions, Compact Blocks costs $2000 \times 5B = 10$ KB, over 3 times the cost of Graphene.

Ordered blocks. Graphene does not specify an order for transactions in the blocks, and instead assumes that transactions are sorted by ID. Bitcoin requires transactions depending on another transaction in the same block to appear later, but a canonical ordering is easy to specify. If a miner would like to order transactions with some proprietary method (e.g., [13]), that ordering would be sent alongside the IBLT. For a block of n items, in the worst case, the list will be $n \log_2(n)$ bits long. Even with this extra data, our approach is much more efficient than Compact Blocks. In terms of the example above, if Graphene was to impose an ordering, the additional cost for $n = 2000$ transactions would be $n \log_2(n)$ bits $= 2000 \times log_2(2000)$ bits $= 2.74$ KB. This increases the cost of Graphene to 5.34 KB, still almost half of Compact Blocks.

3.3.2 Empirical Evaluation

Simulation of Bloom filters and IBLTs. Figure 2 shows the results of a simulation of Compact Blocks and Graphene. The simulation is comprised of many trials, where each trial takes as input a block size (in terms of transactions) and the mempool size. Each protocol is executed and the number of bytes required is recorded. Because

Figure 3.1. A comparison of Graphene and Compact blocks. Mempools are expressed in number of transactions (not bytes) above and beyond the block itself.

Bloom filters and IBLTs are probabilistic mechanisms, the simulation uses the real data structures to ensure the accuracy of the results. The plot is the mean of hundreds of simulations at that point, and error bars are too small to be shown. As the figure shows, Graphene is consistently 1/10 of the cost of Compact Blocks or less, depending on the mempool size. As mempool size increases, the growth of Graphene blocks to Compact Blocks is extremely slow.

Network simulation. Next, we compare Graphene to Compact Blocks with a detailed, custom blockchain simulator using a Python-based discrete event simulator package. Our simulation models the propagation of messages across network links (ignoring effects from variable network bandwidth, TCP, etc.). Nodes accurately model any part of typical blockchain operation necessary for evaluating our metrics, including maintaining a mempool, the blockchain and its forks, and using realistic signaling.

Figure 3.2. When a random graph topology is used, Graphene reduces traffic to 60% of the cost of Compact Blocks.

For Graphene and Compact Blocks, our simulator creates and decodes real Bloom filters and IBLTs, rather than merely estimating whether they might decode or return any false positives. If these data structures fail due to random chance, the nodes recover within the simulation. Because our simulation models detailed signaling and is written in a high-level language, our evaluations are based on a modest number of peers. Since our goal is a comparison between two choices, we expect that our results are representative of larger-scale scenarios.

A challenging parameter to set is the number of transactions per second offered to the network by peers. Our approach is to create kernel density estimates (KDEs) from the transaction generation patterns of real world peers. To that end, we gathered data for all Bitcoin transactions during a three-month period from `http://blockchain.info`. Each transaction in the dataset is labeled with an IP associated with the peer believed to have generated it, as well as the time it was released to the network. For each peer, we normalized the release times by the time of the day in which they were released. We then constructed the KDE for each peer using these normalized transactions times and gaussian kernels with one hour bandwidth. The KDE for a given peer represents a probability distribution from which we can draw transactions over the course of a simulated day. For each peer in the simulator, we randomly select one of the KDEs corresponding to a real world peer. Because these distributions have

been generated from real data, they are a good approximation of the activity of real peers over the average one-day interval. On the other hand, this approach is not able to model days of the week or seasonal phenomena in transaction creation times.

Each simulation is configured to use the following parameters: 1) Topology: a high-degree p2p graph topology. 2) Block Protocol: *Compact Blocks*; or *Graphene*. 3) Block capacity: 2,000 transactions. 4) Full nodes: 50, 100, 150, or 200 peers. In all, we ran 8 combinations of parameters, and we ran each combination with 67 different seeds; all told, we completed 536 simulations. The seeds determined the number of transactions per second (by sampling our KDE, as described above), and the inter-arrival of transactions and blocks. In all simulations, we used 6 miner nodes, representing 6 mining pools. Each simulation was equivalent to 120 minutes; in sum, we simulated about 45 days of blockchain operation. Blocks are generated every 2.5 minutes, like Litecoin; our results would show Graphene to have significantly greater savings if blocks were discovered every 15 seconds (like Ethereum), and show significantly smaller savings if blocks discovered were every 10 minutes (like Bitcoin).

Our main result is shown in Fig. 3.2, where we evaluated the total bandwidth ratio of Graphene to Compact Blocks, as a function of the number of nodes in the network. Since each run is a different number of KBs, we compare the ratio of an exact set of parameters (including the seed), varying only the protocol. Boxplots show the distribution of results across all trials. Fig. 3.2 shows that Graphene reduces traffic to 60% of the cost of using Compact Blocks. As the number of full nodes increases along the x-axis, the ratio of total traffic in the network remains steady, suggesting that our results are representative of larger networks.

CHAPTER 4

RL APPLIED TO BITCOIN

4.1 Background

In this section, we describe double-spend and selfish mining attacks on blockchain systems, and explain the paradigms used to analyze these attacks.

4.1.1 Selfish Mining

The standard Bitcoin protocol requires miners to broadcast a block they mined immediately. However, in the case of *selfish mining*, a miner deliberately withholds transaction information. The motivation is to bifurcate the chain and waste the computational resources of the honest miners should the network decide to build on the attacker's chain. An attacker can't profit economically since the number of blocks that can be created by a miner depends on the fraction of the mining power he has. However, selfish mining discards the honest miners' blocks, by releasing an alternative chain that takes over the current longest chain. If a selfish mining attack is successful, the selfish miners own a higher fraction of the blocks on the main chain because some portion of the blocks created by the honest network go to waste.

Eyal et al. [10] modeled selfish mining using a Markov chain. Then Sapirshtein et al. [18] created a more complex model using a Markov decision process (MDP) for selfish mining and computed the ϵ-optimal policy that increases a selfish miner's revenue. Recently, Gervais et al. [11] incorporated additional parameters such as network conditions and Bitcoin settings into the MDP to study the affects of such parameters on the attacker's policy.

4.1.2 Double Spending

Double-spend attacks [16] work as follows. An attacker creates a transaction that moves funds to a merchant's address. After the transaction appears in the newest block on the main branch, the attacker takes possession of the purchased goods. Using his mining power, the attacker then immediately releases two blocks, with a transaction in the first that moves the funds to a second attacker-owned address. Now the attacker has the goods and his coin back. To defend against the attack, a merchant can refuse to release goods to a customer until z blocks have been added to the blockchain including the first block containing a transaction moving coin to the merchant's address. Nakamoto calculated the probability of the attack succeeding assuming that the miner controlled a given fraction of the mining power [16]; for a given fraction, the probability of success decreases exponentially as z increases.

In general, a merchant may wait z blocks before releasing goods, which can thwart an attacker. But choosing the minimum value of z that secures a transaction is an unresolved issue. The core Bitcoin client shows that a transaction is unconfirmed until it is 6 blocks deep in the blockchain [2], and advice from others is necessarily vague; e.g., "for very large transactions, coin owners might want to wait for a larger number of block confirmations" [4].

Recently, Gervais et al. [11] evaluated the security of blockchains in terms of an attacker's economic profitably and assuming finite resources. They modeled a double-spending attacker's strategy as Markov Decision Process. MDPs are defined by a finite set of discrete states, a set of actions, a transition function, and a reward function. They defined each state in the MDP as a tuple representing the following: the status of the fork, and the number of blocks mined by an attacker, the number of blocks mined by the honest miners, and the number of blocks mined by an *eclipse attack* victim [15], respectively. Gervais et al. encoded several factors into the MDP that affect attacker strategy, including mining power, block depth, connectivity, and the impact of eclipse

attacks. The MDP was implemented with a cutoff value of 20 blocks, representing an attack of finite duration. Using a search algorithm over the space defined by the MDP and a simulation of a blockchain system, they determined the maximum transaction value that would be safe from double spending by an economically rational attacker. This approach is rich, capturing the optimal strategy for double-spending (as well as selfish mining [10, 18]) given network conditions and blockchain parameters. Gervais et al. were able to reach interesting conclusions about the comparative performance and security of several widely used blockchains.

Sapirshtein et al. [18] first observed that some double-spend attacks can be carried out essentially cost-free in the presence of a concurrent selfish mining [10] attack. More recent work extends the scope of double-spends that can benefit from selfish mining to cases where the attacker is capable of *pre-mining* blocks on a secret branch at little or no opportunity cost [20]. The papers identify the optimal mining strategy for an attacker and quantify the advantage he can expect to have over the merchant in terms of pre-mined blocks.

4.2   Problem Statement

The hash rate of miners is the primary quantitative factor that determines the security of any POW based blockchain consensus algorithm. For example, recent MDP models of selfish mining and double-spending assume that the hash rate of honest and malicious miners is known. However, hash rate estimation in blockchain systems can introduce error into these models.

Given target T_i, the expected number of hashes, h, needed to meet the target for a block is

$$\mathbb{E}[h] = \frac{2^{256} - 1}{T_i}.$$

(4.1)

$\mathbb{E}[h]$ describes the *total* number of expected hashes needed to find a block. We have observations regarding the *time* it takes to generate a block. Let $X = X_1, \ldots, X_n$, where $X \sim \text{Exp}(\beta)$ with $\beta = 1/\lambda$. Note that in this scenario, X is a random variable representing the inter-arrival time of blocks discovered by a *single miner*. We could also estimate the network-wide hash rate if we used *all* blocks instead of those by a specific miner. In Section 2.3.2, we have shown that the unbiased MLE estimator for β is equation 2.19. Let r be the hash rate of the network in minutes (or the number of hashes per time unit), and $X = X_1, \ldots, X_n$, where $X \sim \text{Exp}(\beta)$, where $\beta = 1/\lambda$. We can use our estimate, $\hat{\beta}$, of β to calculate a given miner's hash rate.

$$\mathbb{E}[h] = r\lambda = r\frac{1}{\beta} \tag{4.2}$$

$$r = \mathbb{E}[h]\hat{\beta} \tag{4.3}$$

The hash rate computed for a specific miner, given target T_i and n inter-arrival times for $n + 1$ blocks discovered by the miner is

$$\text{Var}(r|T_i, X_1, ..., X_n) = \text{Var}\left(\frac{(2^{256} - 1)\hat{\beta}}{T_i}\right) \tag{4.4}$$

$$= \frac{(2^{256} - 1)^2}{T_i^2}\text{Var}(\hat{\beta}) \tag{4.5}$$

$$= \frac{(2^{256} - 1)^2\beta^2}{T_i^2 n}. \tag{4.6}$$

Example. The current Bitcoin difficulty, D_i, is $1,590,896,927,258 \approx 2^{40}$. Therefore, the current target, T_i, is

$$T_i = \frac{2^{224}}{D_i} = \frac{2^{224}}{2^{40}} \tag{4.7}$$

$$\approx 2^{184}. \tag{4.8}$$

Then the variance associated with the hash rate approximately is

$$\text{Var}(r|T_i, X_1, ..., X_n) = \frac{(2^{256} - 1)^2 \beta^2}{T_i^2 n} \tag{4.9}$$

$$\approx \frac{(2^{256})^2 \beta^2}{T_i^2 n} \tag{4.10}$$

$$\approx \frac{(2^{256})^2 \beta^2}{(2^{184})^2 n} \tag{4.11}$$

$$\approx \frac{2^{512} \beta^2}{2^{368} n} \tag{4.12}$$

$$\approx 2^{144} \frac{\beta^2}{n}. \tag{4.13}$$

The variance associated with estimating hash rate for a miner is large for $0 \le \beta \le 1$ and for any n. We need more blocks than those that are on the Bitcoin blockchain to reduce variance. Reinforcement Learning (RL) provides a solution to this problem because it assumes that an agents acts in an environment where there is uncertainty. Therefore, we use RL methods to search for the optimal strategies for an attacker, who wants to perform selfish mining or double-spending attacks. Then we compare the optimal policy found by RL algorithms to those cited in previous work.

4.3 Preliminary Work

4.3.1 Formulating the Problem as an MDP

States. Using Sapirshtein et al. [18]'s model as a basis, we construct the following MDP that is a 6-tuple $\{S, A, P, R, \gamma, d_0\}$, where $S = \{(x, y, z)\}$ such that $x, y, z \in \mathbb{N}$. The state consists of a 3-tuple where x denotes the number of blocks on the attacker's hidden chain, y denotes the number of blocks created by the honest miners on the main chain, and z denotes the number of blocks that the attacker released on the main chain. Note that this representation assumes that all blocks build on the same parent block.

Actions. $A = \{\texttt{adopt}, \texttt{mine}, \texttt{override}, \texttt{match}\}$. $\texttt{adopt}$ refers to the adoption of the main chain, thereby discarding all blocks created by the attacker. The action $\texttt{mine}$ denotes that the attacker continues to mine, waiting to see who the next block will be discovered by. $\texttt{override}$ refers to an attacker's releasing one more block than the honest miners' blocks on the main chain. This action can be viewed as honest or selfish depending on the current state. If the honest miners have no blocks on the main chain, an addition of a block to the main chain means that the attacker is honest. However, if the honest miners already have blocks on the main chain and the attacker releases an alternative chain that is 1 block longer than that created by the honest miners, then the attacker overwrites the main chain, wasting the victim's computational resources. The $\texttt{match}$ action means that the attacker releases as many blocks as there are on the main chain, causing a bifurcation.

Initial state distribution. $d_0 = \{(0, 0, 0)\}$, where $P(S_0 = (0, 0, 0)) = 1$. In other words, the start state assumes that no blocks have been mined yet. When the attacker chooses the action $\texttt{adopt}$, the agent goes back to the start state because a new parent block is chosen to build on top of.

Transition Function. We consider 3 of parameters of interest included in Gervais et al. [11]'s model. At each time step, a new block is created by the network: with probability q, where q is the mining power of the attacker, the attacker is the winner of a new block. The honest network discovers a block with probability $1 - q$. Not all actions are available in every state. The attacker can always choose the $\texttt{mine}$ action. The $\texttt{adopt}$ action is only available when there are blocks on the main chain created by the honest network. Additionally, $\texttt{override}$ and $\texttt{match}$ are only available when the attackers have enough blocks. At cutoff, $c = 30$, we force the attacker to choose the $\texttt{adopt}$ action in order to have episodic trials. If the number of blocks created by the honest miners exceeds those of the attacker by c, we force the attacker to end the episode. Our third parameter, a, represents network connectivity. If there is a

fork of same length on the main chain, fraction a of the honest miners build on the attacker's alternative chain. We set $a = 1$ to give advantage to the attacker, and to see if RL methods can learn to take advantage of the `match` action. If the attacker matches the main chain with a fork of the same length, all honest miners build on the attacker's chain.

4.4   Proposed Work

CHAPTER 5

TIMELINE

BIBLIOGRAPHY

[1] Andresen, Gavin. O(1) Block Propagation. `https://gist.github.com/gavinandresen/e20c3b5a1d4b97f79ac2`, August 2014.

[2] Confirmation. `https://en.bitcoin.it/wiki/Confirmation`, February 2015.

[3] Bloom, Burton H. Space/Time Trade-offs in Hash Coding with Allowable Errors. *Commun. ACM 13*, 7 (July 1970), 422–426.

[4] Bonneau, Joseph. How long does it take for a bitcoin transaction to be confirmed? `https://coincenter.org/2015/11/what-does-it-mean-for-a-bitcoin-transaction-to-be-confirmed/`, November 2015.

[5] Corallo, Matt. Bip152: Compact block relay. `https://github.com/bitcoin/bips/blob/master/bip-0152.mediawiki`, April 2016.

[6] Difficulty. `http://learnmeabitcoin.com/guide/difficulty`, Feb 2015.

[7] Douceur, J. The Sybil Attack. In *Proc. Intl Wkshp on Peer-to-Peer Systems (IPTPS)* (Mar. 2002).

[8] Eppstein, David, Goodrich, Michael T., Uyeda, Frank, and Varghese, George. What's the Difference?: Efficient Set Reconciliation Without Prior Context. In *ACM SIGCOMM* (2011).

[9] Ethash. `https://github.com/ethereum/wiki/wiki/Ethash`, Last retrieved June 2017.

[10] Eyal, Ittay, and Sirer, Emin Gün. Majority is not enough: Bitcoin mining is vulnerable. *Financial Cryptography* (2014), 436–454.

[11] Gervais, Arthur, O. Karame, Ghassan, Wust, Karl, Glykantzis, Vasileios, Ritzdorf, Hubert, and Capkun, Srdjan. On the Security and Performance of Proof of Work Blockchains. `https://eprint.iacr.org/2016/555`, 2016.

[12] Goodrich, M.T., and Mitzenmacher, M. Invertible bloom lookup tables. In *Conf. on Comm., Control, and Computing* (Sept 2011), pp. 792–799.

[13] Hanke, Timo. A Speedup for Bitcoin Mining. `http://arxiv.org/pdf/1604.00575.pdf` (Rev. 5), March 31 2016.

[14] Hashcash. `https://en.bitcoin.it/wiki/Hashcash`, Last retrieved June 2017.

[15] Heilman, Ethan, Kendler, Alison, Zohar, Aviv, and Goldberg, Sharon. Eclipse Attacks on Bitcoin's Peer-to-peer Network. In *USENIX Security* (2015).

[16] Nakamoto, Satoshi. Bitcoin: A Peer-to-Peer Electronic Cash System, May 2009.

[17] Russel, Rusty. Playing with invertible bloom lookup tables and bitcoin transactions. `http://rustyrussell.github.io/pettycoin/2014/11/05/Playing-with-invertible-bloom-lookup-tables-and-bitcoin-transactions.html`, Nov 2014.

[18] Sapirshtein, Ayelet, Sompolinsky, Yonatan, and Zohar, Aviv. Optimal Selfish Mining Strategies in Bitcoin. `https://arxiv.org/pdf/1507.06183.pdf`, July 2015.

[19] Sompolinsky, Yonatan, and Zohar, Aviv. Secure high-rate transaction processing in Bitcoin. *Financial Cryptography and Data Security* (2015).

[20] Sompolinsky, Yonatan, and Zohar, Aviv. Bitcoin's Security Model Revisited. `https://arxiv.org/abs/1605.09193`, May 2016.

[21] Tschipper, Peter. BUIP010 Xtreme Thinblocks. `https://bitco.in/forum/threads/buip010-passed-xtreme-thinblocks.774/`, Jan 2016.